

Calibrating the Checksum: An Empirical False-Positive Bound for One-Sided ABFT on bfloat16 GEMM, and What It Fails to Catch

Gautam Khosla

<https://github.com/GautamTalksDev/assay-gpu>

August 22, 2026

Abstract

Algorithm-based fault tolerance detects corrupted matrix multiplications by comparing a checksum of the product against a checksum recomputed from the factors, flagging a fault when the residual exceeds a threshold. The classical threshold is $\tau = Ku$, derived from a dot-product rounding bound that requires $Ku \leq 0.01$. At bfloat16 with contraction length $K = 4096$, $Ku \approx 32$, so the bound does not hold and the threshold is not defined. Prior work has noted this and consequently evaluated checksum ABFT only at float32. We supply the missing calibration. This paper is a characterization of a detector under an output-level, post-GEMM perturbation model; it is not a characterization of faults in silicon, and no rate reported here estimates fault incidence in deployed hardware. We fit a Generalized Pareto distribution to 7.99×10^7 clean per-row residuals and extrapolate a threshold at a per-GEMM false-positive rate of 10^{-6} , obtaining 3.1445×10^{-6} , which sits between the two clean sample maxima we observed and is $1.17\times$ the larger of the two runs' floor at $n = 500$. The independence assumption behind the per-GEMM conversion is checked empirically rather than asserted, at a variance ratio of 0.985 against a binomial reference. Against this threshold the detector catches 304 of 400 single exponent-bit flips (76%), failing a bar of 90% that was registered before any data was collected, and catches 0 of 600 low-mantissa flips. Detectability is invariant in K across a $16\times$ range even though the clean noise floor falls as $K^{-0.4999}$: signal and floor scale together, so bit position governs detectability and contraction length does not. This falsified the mechanism the study was designed around. We report four registered predictions that failed, including that one, and four retractions of our own earlier findings, two of which were caught by pre-publication audit rather than by prediction. Code, data, and every pre-registration document with its commit timestamp are released.

1 Introduction

Hardware that returns a wrong answer without saying so is a documented production problem. Google reported silent corrupt execution errors traced to manufacturing defects in CPU cores [Hochschild et al., 2021], Meta reported silent data corruption at fleet scale [Dixit et al., 2021], and both observed rates higher than software engineers expect. The failure mode is not a crash. It is a computation that completes and is wrong.

Matrix multiplication is the natural place to look for this in machine learning workloads, and algorithm-based fault tolerance offers a cheap check. Introduced by Huang and Abraham [1984], ABFT augments a product with checksums and verifies an invariant that holds in exact arithmetic. For $C = AB$ and a vector of ones e , one compares Ce against $A(Be)$: two paths to the same quantity, disagreeing only through rounding, unless something went wrong. The check costs $O(MK + KN)$ against the GEMM's $O(MKN)$.

The difficulty is the threshold. Rounding makes the two paths disagree even when nothing is wrong, so a tolerance is needed, and the classical tolerance comes from a bound on dot-product rounding error that requires $Ku \leq 0.01$, where u is unit roundoff. That condition holds comfortably at float32. It fails at the precisions that machine learning workloads actually use. At an embedding dimension of 4096, $Ku \approx 0.0005$ at float32, $Ku \approx 4$ at float16, and $Ku \approx 32$ at bfloat16 [Ma et al., 2025]. Past float32, the threshold is not conservative; it is undefined.

This has a concrete consequence in the literature. Ma et al. [2025], studying real silent data corruption on unhealthy production nodes, tested whether checksum ABFT could detect it and ran the test at float32 for exactly this reason, reporting that at bfloat16 the check produced false positives on healthy nodes. They concluded that further error analysis is needed to extend ABFT to low-precision datatypes. That is the gap this paper addresses.

We do not derive a better bound. We measure the distribution the bound was supposed to describe, and calibrate a threshold from it. Concretely:

1. We fit a Generalized Pareto distribution to 7.99×10^7 clean per-row residuals at bfloat16, $K = 4096$, and extrapolate a threshold at a stated per-GEMM false-positive rate of 10^{-6} . The method, including the sensitivity checks and the conditions under which we would declare the extrapolation unsound, was fixed in writing before the tail was examined.
2. We check the independence assumption behind the per-row to per-GEMM conversion empirically rather than asserting it.
3. We characterize what the detector catches and misses against that threshold, stratified by which bit of the bfloat16 word was corrupted and by how many bits were corrupted at once.
4. We report that detectability does not vary with contraction length across a $16\times$ range, which falsified the mechanism this study was designed to demonstrate.
5. We evaluate the detector against a detection bar registered before any data existed. It fails.

Two things this paper does not claim. It says nothing about the incidence of faults in real silicon: all corruption here is injected, at the level of the output tensor, under a model specified in Section 3.2. And it claims no novelty for applying ABFT to transformer-shaped GEMMs, which is an active area (Section 2). The contribution is the calibration and the characterization, not the detector.

2 Background and Related Work

2.1 Silent data corruption

Silent data corruption is computation that completes and is wrong. Google traced silent corrupt execution errors to manufacturing defects in CPU cores and reported that the detected incidence is higher than software engineers expect [Hochschild et al., 2021]; Meta reported corruption at fleet scale [Dixit et al., 2021]. Error-correcting codes protect memory and do not address corruption arising inside the datapath, which is what makes detection during computation necessary at all.

2.2 Checksum ABFT

Huang and Abraham [1984] introduced algorithm-based fault tolerance for matrix operations: augment a product with checksums and verify an invariant that holds in exact arithmetic. In the

classical two-sided form, W and X are augmented with a checksum row $e^\top W$ and column Xe . Lighter one-sided variants compute only $e^\top WX$ or WXe , and the matrix sum deviation $e^\top WXe$ reduces the check to a single scalar [surveyed in Xie et al., 2025, Sec. II-C]. The detector studied here is the one-sided form, compared per row rather than reduced to a scalar.

Whatever the form, the detector needs a tolerance, because finite-precision rounding makes the two paths disagree even when the hardware is correct. How that tolerance is chosen is where the literature divides, and it is the axis this paper sits on.

2.3 Thresholds from numerical bounds, and where they stop working

The first approach derives the tolerance from floating-point error analysis. Smith et al. [2015] give the formulation used as a detection rule in subsequent work: report corruption when $\|Cw - A(Bw)\|_\infty > \tau \|A\|_\infty \|B\|_\infty$ with $\tau = Ku$ and u the unit roundoff. The bound behind τ requires $Ku \leq 0.01$ [Golub and Van Loan, 2013, Eq. 2.7.11].

Ma et al. [2025] state the consequence explicitly. For an embedding dimension of 4096: at float32, $Ku \approx 0.0005$ and the condition holds; at float16, $Ku \approx 4$; at bfloat16, $Ku \approx 32$. Past float32 the condition fails, and the threshold is not conservative but undefined. They ran their ABFT evaluation at float32 for this reason, report that running the same check at bfloat16 produced false positives on nodes that had passed production margin tests, and name extending ABFT to low-precision datatypes as open work.

That is the regime this paper occupies: the same contraction length of 4096, at bfloat16, where $Ku \approx 32$ and no principled threshold is available.

2.4 Thresholds from model resilience

The second approach abandons the numerical bound and tunes the tolerance against how much error the model tolerates before its task performance degrades. Xue et al. [2023] relax the strict comparison so that recovery is invoked only when computing errors are large enough to matter, reporting reduced computing overhead and improved accuracy against a strict ABFT baseline. Xie et al. [2025] extend this with a statistical criterion for LLM inference, characterizing resilience across network components and separating error magnitude from error frequency: they observe that components followed by normalization are markedly more sensitive, because a single large error becomes an outlier that skews the mean and variance the normalization depends on. Their thresholds are fitted parameters, set empirically to allow a perplexity increase of 0.3 and an accuracy decrease of 0.5%.

This is calibration, but against a different quantity than ours, and the distinction is not incidental. Both works operate on quantized accelerator datapaths; Xie et al. [2025] inject into INT32 GEMM results with INT8 inputs. In integer arithmetic the checksum identity is exact. There is no rounding, no noise floor, and no false positives to trade against. Their thresholds exist to avoid unnecessary *recovery*, not to avoid spurious *detection*. In floating point the opposite constraint binds: the threshold is bounded below by the rate at which clean hardware trips it, and that rate has to come from somewhere.

2.5 What is missing

Neither family supplies a threshold for floating-point GEMM at a precision where the classical bound has gone vacuous. The numerical-bound line stops at float32 by its own stated condition. The model-resilience line calibrates against task degradation in settings where the false-positive problem does not arise. We are not aware of prior work that calibrates a checksum ABFT threshold

empirically at bfloat16 against a stated per-GEMM false-positive rate, which is what this paper supplies, together with a measurement of what the resulting detector fails to catch.

2.6 Two negative results, from opposite directions

Ma et al. [2025] also report that checksum ABFT failed to flag silent corruption on fourteen of fifteen unhealthy production nodes, and attribute this to real perturbations being small enough relative to the matrix norm to fall inside floating-point error bounds. That is a negative result about the detector, obtained from genuinely faulty hardware, at float32.

This paper reaches a compatible conclusion by a different route: controlled injection at bfloat16, against an empirically calibrated threshold, with the blind spot measured per bit class rather than inferred. The two results share no method, no precision, and no corruption source. They agree that the case for treating checksum ABFT as a complete SDC detector is weak.

2.7 Fault injection and its abstraction level

Fault injection can be performed at the algorithm level, the instruction level, or the micro-architectural level, and the choice bounds what the results support. Instruction-level studies of LLM inference on GPUs exist [Chai et al., 2026]. Algorithm-level injection, used here and in the resilience work above, does not reproduce instruction-level or micro-architectural effects. We treat this as a scope boundary rather than a limitation to be argued away; the model is specified in Section 3.2.

3 Method

This section defines four things: the residual statistic under test, the perturbation model under which it was exercised, the threshold against which detection was scored, and the discipline used to decide what counts as a result. The first three are ordinary experimental description. The fourth matters more than usual here, because most of what follows is a negative result, and a negative result is only worth reading if the bar it failed was set before the data arrived.

3.1 The residual statistic

The detector is a one-sided checksum in the Huang–Abraham family. Let $A \in \mathbb{R}^{M \times K}$ and $B \in \mathbb{R}^{K \times N}$ be the factors, C the computed product, and $e = \mathbf{1}_N$. Two length- M vectors are formed,

$$d = Ce, \quad d' = A(Be), \quad (1)$$

which agree exactly in exact arithmetic. The detector compares them per row, normalized by a nonnegative contraction of the factors:

$$r_i = \frac{|d_i - d'_i|}{s_i}, \quad s = |A|(|B|e), \quad r_{\max} = \max_i r_i. \quad (2)$$

Absolute values are entrywise. All of A , B and C are promoted to `float64` before these operations, and every accumulation in Equation 2 is performed in `float64`. We refer to this statistic as *residual-v3*; the specification is `docs/RESIDUAL.md` in the artifact.

The normalizer is not incidental, and two earlier choices were discarded before arriving at it. Normalizing by the signed grand sum $\max(|\mathbf{1}^\top C \mathbf{1}|, |\mathbf{1}^\top A(B\mathbf{1})|)$ produces a heavy-tailed quotient from two light-tailed pieces, because the denominator measures how completely random signs

cancelled rather than how large the multiply-adds were. On our data that normalizer gave a max-to-median ratio of 389, with Pearson correlation 0.794 between the residual and the reciprocal of its own denominator. The tail was the normalizer, not the hardware.

Replacing the denominator with $e^\top |A||B|e$ removes that failure mode. Every term is nonnegative, so cancellation cannot drive it toward zero, and it concentrates tightly: across $n = 2000$ draws it spanned 1.7169×10^{10} to 1.7193×10^{10} , a max-to-median ratio of 1.0007. It is also the quantity the componentwise theory controls. Writing $C = AB + E$, the checksum functional satisfies $|\mathbf{1}^\top E \mathbf{1}| \leq \gamma_K e^\top |A||B|e$ under the classical bound $|E| \leq \gamma_K |A||B|$ [Higham, 2002]. We use this for its structural content only. At bfloat16 with $K = 4096$ the product Ku exceeds unity and γ_K is not a usable number, and tensor-core paths accumulate in a wider type before storing, so the constant in front of $|A||B|$ is not the classical one either. The claim we rely on is that the checksum functional is bounded by a nonnegative contraction of the factors, and that contraction is the denominator we chose. Every numerical cutoff in this work comes from measurement, not from γ_K .

The final change was to stop reducing before comparing. An earlier version of this work summed the length- M checksum vector to a scalar and compared scalars, discarding $M - 1$ degrees of freedom. Under that formulation, sign flips and mantissa flips were undetectable in every cell we ran, 0 of 1800 samples, and we described the effect as structural. It was not structural. It was the reduction. Under Equation 2, single sign flips are caught in 155 of 200 samples on the same workload, the same shape, the same injection harness and the same seeds. That earlier finding is retracted; Section 4.8 treats it and three others in full.

Two properties of Equation 2 are worth stating because they bound what the detector can be blamed for. The denominator does not depend on C , so a fault in the product cannot shrink the scale that the fault is measured against. And the per-row comparison is only informative if row residuals are not dominated by a shared error term; on the 2000×256 matrix of per-row residuals we measure a mean pairwise Pearson correlation of 0.000126 (standard deviation 0.0225, range -0.0853 to 0.0980), so the rows are effectively independent and the sensitivity gain over the scalar form is not an artifact of correlated noise.

3.2 Perturbation model

Every detection number in this paper was produced under a single, versioned perturbation model, specified in `docs/SPEC-PERTURBATION-MODEL.md` as `perturbation-v1`. We state it early and plainly because it is the strongest constraint on how our results may be read.

After $C = AB$ completes, distinct (element, bit) pairs inside C are XOR-flipped. The residual is then computed from the original factors A and B and the perturbed product. Bits are drawn from one of five classes of the bfloat16 layout, and the number of flips per sample is 1, 2, or 4:

Class	Bits	Width
SIGN	[15]	1
EXPONENT_HIGH	[14:11]	4
EXPONENT_LOW	[10:7]	4
MANTISSA_HIGH	[6:3]	4
MANTISSA_LOW	[2:0]	3

Nothing else is perturbed. Operand A is never flipped, nor is operand B , nor accumulator intermediates inside the GEMM kernel, nor instruction encodings, nor control flow, nor any memory outside the output tensor. This is an exhaustive list, and it has a direct consequence: there is no

operand asymmetry anywhere in our data, because every recorded flip perturbs the numerator of Equation 2 and leaves the denominator untouched.

The obvious objection is that output-tensor XOR is not what a soft error does to a GPU. Instruction-level and micro-architectural effects are invisible to an output-tensor XOR model, and finer-grained injection studies operate below this abstraction [Chai et al., 2026]. We concede the point without qualification. This paper characterizes a detector under a defined perturbation model. It does not characterize faults in silicon, and no rate reported here estimates the incidence of anything in deployed hardware. The fleet silent-corruption literature [Hochschild et al., 2021, Dixit et al., 2021] motivates why a cheap detector would be worth having; it does not supply our numbers, and our numbers do not speak to it.

We verified that the injected perturbations are not silently discarded before they can be measured. Mantissa flips are the case at risk, since a low-mantissa flip followed by a cast could in principle round away. On $n = 50$ samples per class, zero targeted elements were bitwise identical after the cast, and zero samples had a maximum relative delta of zero. The `MANTISSA_HIGH` control produced relative deltas from 0.0317 to 0.492 and `MANTISSA_LOW` from 0.00407 to 0.0313: smaller, as expected for lower bits, but strictly positive on every sample. The undetectability reported in Section 4.4 is a real miss of a real perturbation.

3.3 Workload, hardware, and platform determinism

All measurements use workload W02 in bfloat16 on a single Tesla T4 (sm75), with `torch 2.13.0+cu130`, CUDA 13.0, and driver 580.159.04. Factors are drawn as independent uniform $[-1, 1]$ matrices from a deterministic seed mixer over base seed, workload identifier, shape index, and sample index, with no wall-clock and no hashing, so every draw in this paper is reproducible from its index alone.

Two shape regimes appear. The K-scaling study fixes $M = N = 4096$ and varies $K \in \{512, 1024, 2048, 4096, 8192\}$, with $n = 500$ clean samples per K and $n = 100$ flip samples per $(K, \text{class}, n_{\text{flips}})$ cell. The threshold and false-positive analysis fixes $M = N = K = 4096$, with $n = 2000$ clean samples for the pilot floor, $n = 200$ flip samples per cell, and a separate clean sweep of 20,000 GEMMs described in Section 3.4.

A detector whose clean floor is contaminated by library nondeterminism cannot be characterized at all, so we checked. Across 44 suite cases spanning every GEMM shape in fp32, bf16 and fp16, attention, reductions, transcendentals and a transformer block, all 44 reported bitwise-identical output across repeated launches on this machine. On this hardware, cuBLAS selected a stable kernel every time, which supports treating the clean floor as accumulation error rather than kernel-selection variance. We state the scope of that observation precisely: it is a measured property of one tuple of GPU model, driver, CUDA version and framework build. It is not a general property of cuBLAS and must not be carried onto other hardware.

3.4 Threshold and detection predicate

Residual-v3 performs $M = 4096$ per-row tests per GEMM. A per-GEMM false-positive rate below 10^{-6} , which is the bar this detector was built against, therefore requires a per-row rate of $1 - (1 - 10^{-6})^{1/4096} \approx 2.44 \times 10^{-10}$ under independence. No sampling budget we could afford estimates that directly, so it has to be extrapolated, and the extrapolation method was fixed in writing before the tail was examined.

We collect clean per-row residuals from 20,000 GEMMs at $K = 4096$, using the full 4096-row vector rather than a subsample, for roughly 8×10^7 row observations. Thresholds at the 95th, 99th and 99.9th percentiles are frozen from an initial pass of 500 GEMMs, giving 1.057096×10^{-6} ,

1.392735×10^{-6} and 1.776693×10^{-6} . A Generalized Pareto distribution is fitted to exceedances above each frozen threshold by maximum likelihood, chosen because it is the standard peaks-over-threshold model and not because it fits these data. The $1 - 2.44 \times 10^{-10}$ quantile of the fitted distribution is the threshold. Fitting at three separate percentiles is a pre-registered sensitivity check, and the pre-registered rule was that thresholds differing by more than a factor of two are to be reported as untrustworthy rather than averaged.

The canonical threshold is the extrapolation from the 99th-percentile fit,

$$\tau = 3.144535556 \times 10^{-6}, \quad (3)$$

and every detection rate reported in this paper is scored against it unless explicitly stated otherwise. Section 4.5 reports the fitted parameters, the sensitivity comparison, and the one threshold that turned out not to be evaluable.

Two points of provenance need stating here rather than in a limitations section. First, τ is not the largest clean residual we observed. Observed maxima are sample maxima, they grow with n , and two of them appear in our own records for the same configuration at different sample sizes: 3.245921×10^{-6} at $n = 2000$ and 2.684525×10^{-6} at $n = 500$. Detection rates differ by a few points depending on which is used, with no change to any rule or fit. We report τ as canonical because it is the only threshold with a stated per-GEMM false-positive rate behind it, and we note that it falls between the two sample maxima, so it is not an artifact of either sampling run. Second, the conversion from per-GEMM to per-row rate assumes independence across rows. Rather than assert it, we test it: the number of per-row exceedances above the 99th percentile has observed mean 40.42 and variance 39.93 per GEMM, against 40.96 and 40.55 for $\text{Binomial}(4096, 0.01)$, a variance ratio of 0.985. The counts are very slightly under-dispersed relative to independent Bernoulli trials. This supports the conversion at the fitting threshold, and it does not establish independence deep in the tail, which we carry forward as a stated assumption.

The detection predicate itself requires one sentence that would otherwise seem too obvious to write. A sample is detected when $r_{\max} > \tau$ or $r_{\max}$ is not finite. An infinite residual means the checksum diverged, which is the loudest signal the detector can emit, and a naive comparison against a threshold silently evaluates false on it. We found this in our own analysis code during a pre-publication audit, where it had understated detection in ten of fifteen `EXPONENT_HIGH` cells by as much as fourteen percentage points and manufactured an apparent effect that does not exist. Section 4.8 records the correction. We note the general form of the hazard here because it applies to any threshold detector operating on floating-point residuals: the case that most obviously indicates failure is the case that comparison operators handle worst.

3.5 Pre-registration

Every headline claim in this paper was predicted before it was measured. Predictions were committed to the repository, with numeric ranges and an explicit falsifier, before the corresponding data existed, and the commit history establishes the ordering independently of our word for it. Four prediction sets appear: the noise-floor and sensitivity predictions for residual-v3, the K-scaling predictions, the extreme-value method and its predictions, and a later mechanism prediction about operand asymmetry.

The rule we imposed on ourselves is that a prediction may be falsified by new data but may not be rewritten by it. That rule was tested repeatedly and is the reason this paper is worth reading rather than the reason it is short. Of the K-scaling predictions, three of five failed, including the one that motivated the experiment. The pre-registered falsifier for the underlying mechanism fired. In the residual-v3 set, two of five predictions missed, one of them by two percentage points, and

K	median r_{median}	observed clean max
512	1.029719×10^{-6}	7.636551×10^{-6}
1024	7.281700×10^{-7}	5.772809×10^{-6}
2048	5.150956×10^{-7}	3.746226×10^{-6}
4096	3.640184×10^{-7}	2.684525×10^{-6}
8192	2.575653×10^{-7}	1.940381×10^{-6}

Table 1: Clean noise floor against contraction length. $n = 500$ per K .

it is reported as a miss rather than rounded to a pass. A mechanism prediction registered late in the project was falsified at its precondition when we discovered the harness could not test it, and it is reported as such rather than deleted. The kill-test bar that this detector failed is the same number that appeared in the repository README before any code was written, and it has never been moved.

We also state where pre-registration did not save us. Two defects reached published internal documents and were caught by audit rather than by prediction. The first is the detection predicate described above, withdrawn in Section 4.8. The second was a results document citing a source data file under a path that never existed; the file was located, the numbers it produced were reproduced from it, and the provenance was corrected in the repository. Writing the bar down first constrains what you may conclude from a measurement. It does not guarantee the measurement is right.

4 Results

All detection rates in this section are scored against $\tau = 3.144535556 \times 10^{-6}$ unless stated otherwise, using the predicate of Section 3.4: detected when $r_{\text{max}} > \tau$ or r_{max} is not finite.

4.1 The clean noise floor scales as $K^{-1/2}$

Across 2500 clean draws, 500 at each of five contraction lengths with $M = N = 4096$ fixed, the median per-row residual falls monotonically with K (Table 1). A log-log fit over the full $16\times$ range gives an exponent of -0.4999 ; restricted fits give -0.4997 over 512 to 2048 and -0.4991 over 4096 to 8192.

The mechanism is a ratio. Absolute accumulation error in a row sum of K bfloat16 products grows as $\sqrt{K}$, while the per-row normalizer $s_i = |A_i||B|e$ grows linearly in K . The relative residual therefore falls as $\sqrt{K}/K = K^{-1/2}$.

We record a correction here. Our locked prediction stated the mechanism as “clean accumulation error grows as $\sqrt{K}$, signal-to-floor should fall as $1/\sqrt{K}$.” The measured exponent matches $\sqrt{K}$ scaling in magnitude, but that reasoning conflated absolute and relative error and predicted the wrong direction for the floor’s dependence on K . The exponent was right for a reason the argument did not supply.

4.2 Detectability does not vary with K

The study was designed to demonstrate that detection degrades as contraction length grows: a fixed single-element perturbation should become harder to see as the row sum it hides in gets larger. We registered that prediction, together with a falsifier stating that if detection were flat across a

bit class	n_{flips}	@ clean max	@ τ	Δ (pp)
EXPONENT_HIGH	1	168/200	166/200	-1.00
EXPONENT_HIGH	2	194/200	193/200	-0.50
EXPONENT_HIGH	4	200/200	200/200	+0.00
EXPONENT_LOW	1	152/200	138/200	-7.00
EXPONENT_LOW	2	184/200	179/200	-2.50
EXPONENT_LOW	4	199/200	196/200	-1.50
MANTISSA_HIGH	1	14/200	10/200	-2.00
MANTISSA_HIGH	2	22/200	17/200	-2.50
MANTISSA_HIGH	4	32/200	24/200	-4.00
MANTISSA_LOW	1	1/200	0/200	-0.50
MANTISSA_LOW	2	1/200	0/200	-0.50
MANTISSA_LOW	4	1/200	0/200	-0.50
SIGN	1	162/200	156/200	-3.00
SIGN	2	194/200	189/200	-2.50
SIGN	4	200/200	200/200	+0.00

Table 2: Detection by bit class and flip count at $K = 4096$, $n = 200$ per cell. “Clean max” is 2.684525×10^{-6} ; τ is $3.144535556 \times 10^{-6}$. No cell detects better at the stricter threshold.

$16\times$ range of K , the mechanism would be wrong and the finding would reduce to a single-shape observation with no predictive content.

The falsifier fired. Single sign-bit flips are detected at 79%, 75%, 73%, 81%, and 77% at $K = 512$ through 8192: a mean of 77% and a spread of eight percentage points across a $16\times$ range, with no monotone trend and variation consistent with binomial noise at $n = 100$ (standard error roughly four points).

The explanation is in Section 4.1. The threshold is derived from the clean floor at each K , and the clean floor falls as $K^{-1/2}$. A perturbation of fixed magnitude produces a residual that also scales against the same growing normalizer. Signal and floor move together, so their ratio, which is what the detector tests, does not depend on K . Contraction length does not govern detectability; bit position does (Section 4.3).

Threshold estimator consistency. Detection in this study uses the per- K observed clean maximum as the threshold, taken over 500 clean samples at that K . An observed maximum is a sample statistic that grows with n and is therefore not a stable absolute reference; this is precisely why τ exists. It does, however, use the same estimator at the same sample size at every K , so the cross- K comparison is internally consistent even though each individual threshold is unstable in absolute terms. We state this as a limitation with a defense, not as a claim that these thresholds are stable.

4.3 Bit position governs detectability

Table 2 gives the full flip matrix at $K = 4096$, $n = 200$ per cell, scored at τ and at the observed clean maximum of the K-scaling run for comparison.

The stratification spans the full range. Exponent flips are caught most of the time, sign flips slightly less often, high mantissa flips rarely, and low mantissa flips never. This ordering is stable across every contraction length we tested. It is a statement about perturbation magnitude relative

fit	ξ	95% CI	extrapolated τ
p99, exclude truncated	-0.07806	$[-0.08000, -0.07610]$	3.144536×10^{-6}
p99, right censored	-0.07809	$[-0.08003, -0.07613]$	3.144314×10^{-6}
p99.9	-0.05318	$[-0.05960, -0.04661]$	3.341952×10^{-6}
p95	NOT EVALUABLE		

Table 3: Generalized Pareto fits to clean per-row residual exceedances, extrapolated to the $1 - 2.44 \times 10^{-10}$ quantile. Fits are reported separately and not averaged.

to the accumulation floor, expressed through which field of the bfloat16 word was disturbed, and it does not depend on K .

4.4 Low mantissa flips are a hard negative

MANTISSA_LOW is detected in zero samples in every cell we ran. Two counts, which must not be merged because they rest on different thresholds: 0/1500 across the K-scaling study scored at per- K observed clean maxima, and 0/600 at τ with $K = 4096$.

At $K = 4096$ the minimum, median, and maximum ratio are identical to four significant figures at one, two, and four flips (0.5247, 0.6217, 0.827; these ratios are to the pilot clean maximum of 3.245921×10^{-6} , the basis on which they were originally recorded). More corruption does not produce a larger residual. A mantissa-LSB perturbation lies below the accumulation floor at this contraction length regardless of how many of them there are.

We tested whether this was an artifact of the injector rather than a property of the detector, on the hypothesis that a low-mantissa flip might be rounded away by the cast back to the working dtype. It is not. Across 50 samples per class, zero targeted elements were bitwise identical after the cast and zero samples had a maximum relative delta of zero; the MANTISSA_HIGH control produced relative deltas from 0.0317 to 0.492 and MANTISSA_LOW from 0.00407 to 0.0313. The flips are real and the detector does not see them.

The shared minimum of 0.5247 across all six mantissa cells is unexplained. Section 4.8 describes a registered prediction that would have explained it and could not be tested.

4.5 A calibrated false-positive bound

Table 3 reports the fits. The shape parameter is negative at both evaluable thresholds with confidence intervals excluding zero, indicating a bounded upper tail, and the extrapolated thresholds are $1.17\times$ and $1.24\times$ the observed clean maximum.

The pre-registered sensitivity check required the refits to agree within a factor of two, failing which the extrapolation was to be reported as untrustworthy rather than averaged. The two evaluable fits agree at $1.063\times$. The third could not be evaluated, for reasons in Section 5.

The per-row target of 2.44×10^{-10} follows from a per-GEMM target of 10^{-6} only under independence across the 4096 rows. Rather than assume this, we measured the dispersion of per-row exceedance counts: observed mean 40.42 and variance 39.93 per GEMM, against 40.96 and 40.55 for $\text{Binomial}(4096, 0.01)$, a variance ratio of 0.985. The counts are slightly under-dispersed relative to independent Bernoulli trials, which supports the conversion at the fitting threshold. It does not establish independence in the far tail, and we carry that forward as a stated assumption. Separately, the mean pairwise correlation between per-row residuals is 0.000126.

ID	Claim	Measured	Verdict
Q1	floor exponent in $[0.4, 0.6]$	0.4999	PASS
Q2	sign detection falls monotonically with K	non-monotone	FAIL
Q3	sign detection $> 95\%$ at $K = 512$	79%	FAIL
Q4	sign detection $< 70\%$ at $K = 8192$	77%	FAIL
Q5	MANTISSA_LOW 0/100 at every K	0/1500	PASS
P1	clean per-row floor median $2-8 \times 10^{-6}$	3.64×10^{-7}	MISS
P2	clean per-row max/median 4-15	8.92	PASS
P3	sign 1-flip median ratio 2-10	2.298	PASS
P4	MANTISSA_LOW 1-flip median ratio < 1	0.6217 [†]	PASS
P5	EXPONENT_LOW 1-flip detection $> 70\%$	68%	MISS
F1	GPD ξ negative, in $[-0.5, 0]$	$-0.078, -0.053$	PASS
F2	threshold $1-5\times$ clean max	$1.17\times, 1.24\times$	PASS
F3	three refits agree within $2\times$	two agree at $1.06\times$	PARTIAL
F4	sign detection falls < 10 pp at τ	-3 pp	PASS

Table 4: Registered predictions and outcomes. Q and P sets were locked before their data existed; the F set was locked before the residual tail was examined. [†]ratio to the pilot clean maximum 3.245921×10^{-6} .

4.6 The detector fails its registered bar

The bar was written into the repository README before any code was written: at least 90% of exponent-bit flips detected, at a per-GEMM false-positive rate below 10^{-6} . Pooling the two exponent classes at a single flip, the detector catches 304 of 400, or 76%, at τ . That is a failure, and the bar was never moved.

Two points of provenance. First, the same pooled quantity is 301/400 (75.25%) at the pilot observed clean maximum of 3.245921×10^{-6} and 320/400 (80%) at the K-scaling clean maximum of 2.684525×10^{-6} . These differ because an observed maximum is a sample statistic that grows with n , not because any rule or fit changed. We report 76% as canonical because τ is the only threshold with a stated false-positive rate behind it, and we note that τ falls between the two sample maxima and is therefore not an artifact of either run.

Second, the verdict was determined before the fit existed. We recorded, in advance, that the extrapolated threshold was expected to be at or above the observed clean maximum and that detection is monotone non-increasing in threshold, so 75.25% bounded detection from above at any defensible threshold and the bar was already unreachable. The measurement went the predicted direction: detection fell rather than rose.

4.7 Registered predictions that failed

Table 4 gives every registered prediction and its outcome. All were committed before the corresponding data existed.

P5 is recorded as a miss at $136/200 = 68\%$ against a bar of 70%. It is two points short and we have not rounded it up. Q2 through Q4 are the mechanism this study was built on, and all three failed.

4.8 Retractions

Four findings in this project were withdrawn. We report them because a paper whose principal result is a negative one is only worth reading if the process that produced it is visible.

One: the scalar reduction. An earlier residual summed the length- M checksum vector to a scalar before comparing. Under it, sign flips and mantissa flips were undetectable in every cell, 0 of 1800 samples, and pooled exponent detection at one flip was $179/400 = 44.75\%$ with an exact binomial 95% interval of $[39.8\%, 49.8\%]$. We described the sign and mantissa result as structural invisibility of the ones-vector checksum. It was not structural. It was the reduction discarding $M-1$ degrees of freedom before the comparison. Under the per-row form, sign detection is $155/200$ on the same workload, shape, harness, and seeds. The earlier evaluation is void and retained in the repository as evidence that the bar was not moved.

Two: the mantissa near-miss. We suspected the MANTISSA_LOW zero was an injector artifact and tested it. It was not (Section 4.4). The hypothesis is withdrawn and the hard negative stands.

Three: a prediction falsified at its precondition. We registered a mechanism to explain the MANTISSA_LOW floor and the shared 0.5247 minimum: a flip in operand A perturbs both the numerator and the denominator of Equation 2 and may self-cancel, while a flip in B perturbs the numerator only, so detectability should differ by operand. We committed four predictions with numeric separations. We then found that the harness injects into the output tensor only, so no operand asymmetry exists in any of our data and the predictions are untestable rather than merely unmeasured. No pass or fail is claimed for them. The 0.5247 minimum remains unexplained. Operand-level injection is future work.

Four: a non-monotonicity that was a comparison bug. We reported that EXPONENT_HIGH detection was non-monotone in flip count, observing 91% at two flips and 85% at four at $K = 1024$, with a similar inversion at $K = 4096$ and an apparently independent occurrence in a separate run. Pre-publication audit found the cause. Exponent flips can produce a non-finite residual, and the analysis path scored detection with a bare comparison, which evaluates false on a non-finite value. Ten of fifteen EXPONENT_HIGH cells contained non-finite residuals, up to 14 of 100 at $K = 1024$ with four flips; no other bit class produced any. Under the correct predicate, detection is monotone at every K : 78/97/100 at $K = 512$, 80/96/99 at 1024, 88/97/100 at 2048, 83/97/100 at 4096, and 83/97/100 at 8192. The reported effect does not exist and is withdrawn. The corroborating run was the same defect, not independent confirmation.

This last one is worth generalizing. An infinite residual is the loudest signal this detector can emit, and a naive threshold comparison silently discards it. The hazard applies to any threshold detector operating on floating-point residuals.

5 Scope

These are boundaries on what was measured, not a list of things we intend to fix.

Perturbation level. All injection is at the output tensor, post-GEMM. Operands, accumulator intermediates, instruction encoding, and control flow are never perturbed. Nothing here estimates fault incidence in silicon.

Single configuration. One GPU model (Tesla T4, sm75), one workload, one datatype (bfloat16), one shape family with $M = N = 4096$. The clean floor is per-configuration by construction and must not be transferred to other hardware.

The p95 fit is not evaluable. Our data collection stored the top 64 per-row residuals per GEMM. At the 95th percentile the mean number of exceedances is 205.3 per GEMM with a maximum of 263, so all 19500 GEMMs are truncated and roughly two thirds of exceedances are lost, systematically the smaller ones. At the 99th percentile only 2 of 19500 GEMMs exceed the capture and the effect is negligible: the shape parameter differs between excluding and right-censoring the truncated samples by 2.59×10^{-5} against a confidence interval width of 3.90×10^{-3} , or 0.66% of the interval. This is a data-collection limitation discovered after the method was locked. It is not a change of method and not a redefinition of the sensitivity check, which remains as registered with two of three refits evaluable.

Extrapolation distance. The fit reaches from an observed exceedance probability of order 10^{-2} to 2.44×10^{-10} , several decades beyond the data. The implied upper endpoints of the fitted distributions differ more than the reported quantiles do: 3.74×10^{-6} from the p99 fit against 4.60×10^{-6} from p99.9, a spread of 1.23 $\times$, versus 1.06 $\times$ agreement at the extrapolated quantile. The hard ceiling on the residual is less well determined than the sensitivity ratio alone suggests.

Sample sizes. Flip cells are $n = 100$ in the K-scaling study and $n = 200$ at $K = 4096$, giving standard errors of roughly four and three percentage points. Sub-percent detection probabilities are not resolved.

Unexplained observations. The shared minimum ratio of 0.5247 across all mantissa cells is unexplained and is reported as an open observation.

Known open issue. Two checksum reduction backends in our implementation disagree, one accumulating in float32 and one in float64. All results reported here use a single path with float64 accumulation throughout, but the discrepancy is unresolved and a threshold calibrated on one backend should not be assumed to transfer to the other.

6 Discussion

The headline is a failure, and the useful part is the shape of it.

The detector is close to reliable on exponent corruption and blind to low mantissa corruption, and no amount of additional corruption in the low mantissa changes that. This is not a tuning problem. A perturbation that lies below the accumulation floor is not recoverable by lowering the threshold, because the threshold is bounded below by the false-positive requirement. The 76% figure is what this detector achieves at a defensible operating point, not what it achieves at a favorable one.

Calibrating empirically also changes what the threshold means. The classical $\tau = Ku$ is a bound: below it, corruption is guaranteed invisible. Ours is a quantile: above it, false positives

occur at a stated rate. The second is weaker as a guarantee and usable in a regime where the first is undefined. A practitioner deploying checksum ABFT at bfloat16 today has no principled threshold available, and the choice is between a calibrated quantile and no check at all.

The K-invariance result is the one we would flag for anyone extending this. The intuition that longer contractions hide faults better is natural, we held it firmly enough to design a study around it, and it is wrong for this detector, because the normalizer grows with K at the same rate the noise does. Anyone reasoning about ABFT sensitivity from row-sum magnitude alone should check whether their normalizer cancels the effect they are predicting.

Finally, the agreement with Ma et al. [2025] is worth stating plainly. They found checksum ABFT failing to catch real corruption on real unhealthy nodes at float32. We find it catching roughly three quarters of single exponent flips and none of the low mantissa flips at bfloat16 under controlled injection with a calibrated threshold. Different corruption sources, different precisions, different methods, compatible conclusions. The case for treating checksum ABFT as a complete SDC detector is weak, and the case for using it as a cheap first filter with known blind spots is what the evidence supports.

7 Conclusion

We calibrated a detection threshold for one-sided checksum ABFT at bfloat16 and $K = 4096$, a regime where the classical threshold is undefined, and characterized what the detector catches and misses against it. It catches 76% of single exponent-bit flips against a registered bar of 90%, and none of 600 low-mantissa flips. Detectability is governed by bit position and not by contraction length, which falsified the mechanism the study was designed around.

Four registered predictions failed and four findings were retracted, two of them caught by audit rather than by prediction. We report all of it. The threshold, the perturbation model, the predictions with their commit timestamps, and the data are released so that each number here can be checked against the file that produced it.

Artifact Availability

Code, data, specifications, and pre-registration documents are available under Apache 2.0 at <https://github.com/GautamTalksDev/assay-gpu>, at tag v0.1.0-paper. The dataset is released under CC BY 4.0.

References

- Duo Chai, Zizhen Liu, Shuhuai Wang, Songwei Pei, Cheng Liu, Huawei Li, and Shangguang Wang. Analysis of LLM vulnerability to GPU soft errors: An instruction-level fault injection study, 2026.
- Harish Dattatraya Dixit, Sneha Pendharkar, Matt Beadon, Chris Mason, Tejasvi Chakravarthy, Bharath Muthiah, and Sriram Sankar. Silent data corruptions at scale, 2021.
- Gene H. Golub and Charles F. Van Loan. *Matrix Computations*. The Johns Hopkins University Press, Baltimore, MD, 4 edition, 2013. Eq. 2.7.11, dot-product rounding bound.
- Nicholas J. Higham. *Accuracy and Stability of Numerical Algorithms*. Society for Industrial and Applied Mathematics, Philadelphia, PA, 2 edition, 2002. doi: 10.1137/1.9780898718027.

- Peter H. Hochschild, Paul Turner, Jeffrey C. Mogul, Rama Govindaraju, Parthasarathy Ranganathan, David E. Culler, and Amin Vahdat. Cores that don't count. In *Proceedings of the Workshop on Hot Topics in Operating Systems (HotOS '21)*, pages 9–16, Ann Arbor, MI, USA, 2021. ACM. doi: 10.1145/3458336.3465297.
- Kuang-Hua Huang and Jacob A. Abraham. Algorithm-based fault tolerance for matrix operations. *IEEE Transactions on Computers*, C-33(6):518–528, 1984. doi: 10.1109/TC.1984.1676475.
- Jeffrey Ma, Hengzhi Pei, Leonard Lausen, and George Karypis. Understanding silent data corruption in LLM training, 2025. Appendix D evaluates checksum ABFT and reports the low-precision constraint.
- Tyler M. Smith, Robert A. van de Geijn, Mikhail Smelyanskiy, and Enrique S. Quintana-Ortí. Towards ABFT for BLIS GEMM. Technical Report FLAME Working Note #76, TR-15-05, The University of Texas at Austin, June 2015. URL <https://www.cs.utexas.edu/~flame/pubs/FLAWN76.pdf>. Revised 5 November 2015.
- Tong Xie, Jiawang Zhao, Zishen Wan, Zuodong Zhang, Yuan Wang, Runsheng Wang, Ru Huang, and Meng Li. ReaLM: Reliable and efficient large language model inference with statistical algorithm-based fault tolerance. In *Proceedings of the 62nd ACM/IEEE Design Automation Conference (DAC)*, 2025.
- Xinghua Xue, Cheng Liu, Haitong Huang, Ying Wang, Bingru Yang, Tao Luo, Lei Zhang, Huawei Li, and Xiaowei Li. ApproxABFT: Approximate algorithm-based fault tolerance for vision transformers, 2023.